Truthprint: Typed Invariant Contracts for Authenticated Semantic Provenance Across Machine Translation

Anonymous ACL submission

Abstract

Token-level LLM watermarks are fragile under translation and paraphrasing because their signal lives in tokens. We propose *Truthprint*, a semantic intermediate representation that separates *locked invariants* (polarity, quantity, time, attribution, roles) from *controlled realization* via a machine-checkable typed mutability contract. What survives transformation is *ledger-based authenticated semantic provenance*—a keyed tag over the invariant digest, re-verified after transformation—not an in-text watermark: hiding a payload in realization freedom is a negative result under real MT (a surface parser reads 1 of 4,800 NLLB outputs). On 4,800 NLLB translations across six languages, a Stage-2 extractor recovers the invariants (0.83–0.99) and the scheme attributes translated documents cross-lingually (core6 document attribution 0.90–1.00), where a token watermark (KGW) and a SemStamp-style embedding-region detector largely fail, and rejects single-field tampers of contract fields (≥ 0.996), which robust hashing does not. Benign rule-based paraphrase authenticates (1.00); a fixed-lexicon adaptive attacker removes the mark, while two small open-vocabulary neural runs resist a held-out attack (ValidRemoval 0.05/0.10, N=40/10). We state two standard reductions to the MAC and show that the operative false-positive floor is the contract-collision rate, not $2^{-\tau}$. All results are closed-domain and small-scale; Truthprint verifies consistency with a participating keyed generator, not arbitrary AI text.

Keywords. Intermediate representation, authenticated provenance, LLM watermarking, semantic provenance, translation robustness

1 Introduction

Provenance for LLM text must survive ordinary transformation. Generation-time watermarking biases decoding so that an authorized detector finds a statistical signature; SynthID-Text integrates this into sampling without retraining and detects it without rerunning the generator (Dathathri et al., 2024).

Article 50 of the EU AI Act motivates machine-readable marking (European Parliament and Council, 2024), useful only if robust (Liu et al., 2024c); robustness here is a technical objective, not a compliance claim. Provenance also aids data hygiene, since training on unattributed synthetic text degrades models over generations (Shumailov et al., 2024).

Token-level methods embed their signal in surface token choices, which translation, paraphrasing, and summarization replace while preserving meaning: the provenance claim concerns content identity, but the carrier depends on language-specific lexical decisions.

Recent semantic watermarks use sentence embeddings, contextual representations, or structured graphs (Hou et al., 2024; Ren et al., 2024; Liu et al., 2024b; Ye et al., 2026; Schäfer et al., 2026), but semantic similarity alone does not distinguish meaning that must never change from choices that may safely vary: two sentences can be embedding-close yet differ in negation, quantity, temporal scope, attribution, modality, or causal direction.

This paper proposes *Truthprint*, a typed semantic IR for provenance that explicitly separates:

1. **semantic invariants**, which must remain fixed; and
2. **controlled realization variables**, which may encode watermark symbols only when

semantic and pragmatic equivalence is validated.

The central novelty is a machine-checkable contract that fixes which semantic fields are locked and which realization choices may vary; a keyed tag over the locked fields authenticates meaning out of band—the mechanism that survives real translation, unlike an in-text watermark constrained by the same contract (Sec. 10.1). Figure 1 shows the end-to-end data flow.

The delta over the closest prior work is precise. Embedding-based semantic watermarks (Hou et al., 2024; Ren et al., 2024; Liu et al., 2024b; Schäfer et al., 2026) define admissibility through a *metric* (closeness under a learned distance); Truthprint defines it through a *typed predicate*—admissible only when a symbolic invariant relation holds field by field—so a watermark symbol ranges over key-dependent invariant-equivalence classes rather than tokens or embedding regions (Sec. 9). Conditional on sound extraction, encoding stays within a class and preserves protected fields, and because attribution is gated by a payload bound to the invariant digest, fresh-payload forgery reduces to MAC security. SWAN (Ye et al., 2026) shares a graph carrier but adopts an existing formalism wholesale, whereas Truthprint types each field with an explicit mutability policy validated before release.

1.1 Contributions

- **Method.** We formulate watermark carriers as secret-keyed realization choices constrained by a field-level mutability contract (*locked*/*conditional*/*realization-only*), combining uncertain-carrier erasures, error correction, and invariant-bound payload authentication.

- **Conditional analysis.** Assuming sound invariant extraction and validation we establish protected-field fidelity; under standard cryptographic assumptions fresh-payload forgery and per-test tag acceptance reduce to the MAC; and we show the operative false-positive floor is the contract-collision rate, with an analytic crossover to the $2^{-\tau}$ bound.

- **Reproducible real-MT evidence.** On 4,800 released NLLB translations across six languages the meaning layer survives and authenticates cross-lingually where surface carriers and a real token watermark do not; at a matched FPR it survives translation at least as well as a SemStamp-style detector and robust hashing (a strong encoder lets SimHash match it) and alone rejects single-field tampers; searched without a document id, it stays specific up to $\sim 10^3$ registered documents but not at 10^5; and two small open-vocabulary neural runs resist a held-out adaptive attack. All Truthprint numbers reproduce from the released data with no GPU; wide-coverage parsing and scale remain open.

Claim–evidence boundary. Table 7 (appendix) separates what the artifact *demonstrates*, what is proved only under explicit assumptions, and what remains a hypothesis (e.g., "language-neutral" names the intended schema interface, not a measured property).

2 Background and Related Work

Token-level and provable watermarks. The green–red list watermark (Kirchenbauer et al., 2023) and SynthID-Text (Dathathri et al., 2024; Google DeepMind, 2024; Google, 2026), which generalizes it via tournament sampling, embed a keyed statistic in token selection; analysis shows detector robustness depends on its statistical assumptions (Omidi et al., 2026). A parallel line seeks formal guarantees: provable robustness for a unigram watermark (Zhao et al., 2024), distortion-free editing-robust watermarks (Kuditipudi et al., 2024), cryptographically undetectable watermarks (Christ et al., 2024), publicly verifiable unforgeable watermarks (Liu et al., 2024a), and low-FPR statistical tests (Fernandez et al., 2023). All carry the signal in token decisions, which meaning-preserving transformations replace; He et al. show token watermarks lack cross-lingual consistency and are removed by pivot translation (He et al., 2024). Truthprint inherits the cryptographic spirit at the *provenance* layer—fresh-payload unforgeability and per-candidate tag acceptance (Propositions 1, 2) reduce to MAC security—with no claim of undetectability.

Semantic and structured watermarks. Embedding-region and semantic-context

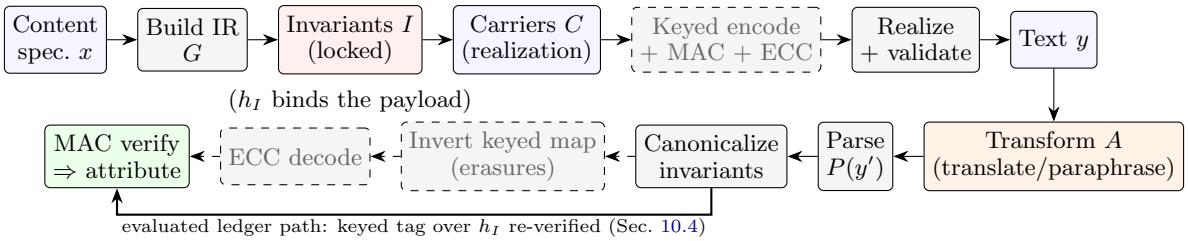


Figure 1: Truthprint data flow. Encoding (top) locks the invariant set I and validates every candidate realization against the contract before emitting text; detection (bottom) re-parses possibly transformed text and canonicalizes the invariants. Dashed boxes form the in-text carrier path (keyed encoding, inversion, ECC), a negative result under real MT (Sec. 10.1). The evaluated mechanism is the solid bypass: a keyed tag over the invariant digest h_I, registered out of band and re-verified after re-extraction (Sec. 10.4); an invariant-altering transform breaks it.

methods—SemStamp (Hou et al., 2024), SemaMark (Ren et al., 2024), semantic-invariant watermarking (Liu et al., 2024b), dual embeddings (Schäfer et al., 2026)—reduce token dependence; SWAN carries the signal in AMR structure (Ye et al., 2026), and SynGuard hybridizes token masking with semantic alignment (Han et al., 2025). But embedding geometry does not expose *which* typed fact changed: high similarity can coexist with altered polarity, quantity, attribution, modality, or causal direction. Truthprint instead defines a watermark-specific typed IR, assigns each field a mutability policy, and validates every realization against an immutable contract before release.

Attacks and robust hashing. Watermark presence and family can be identified by black-box probing (Gloaguen et al., 2025), motivating a Kerckhoffs threat model (security in the key, not obscurity); scrubbing and spoofing under distillation (Yi et al., 2025) and radioactive traces in fine-tuned models (Sander et al., 2024) motivate MAC-authenticated payloads rather than statistical agreement; and detector-guided adversarial paraphrasing (Cheng et al., 2025) motivates counting a removal as successful only when it both evades detection and preserves meaning:

$$\text{ValidRemoval} = \mathbf{1}\big[D_K(y') < \tau_D\big] \cdot \mathbf{1}\big[\text{InvariantEq}(P(y), P(y')) = 1\big]. \quad (1)$$

Retrieval over stored generations already defends against paraphrase (Krishna et al., 2023); Truthprint stores a typed record instead of the text, adding field-level tamper rejection (Sec. 10.6). Registering a meaning-digest tag and re-verifying it after transformation also resembles *robust hashing* (MinHash (Broder, 1997), SimHash (Charikar, 2002), learned semantic hashing), but those answer a *similarity* question that cannot detect a typed field change—a single-field edit stays within tolerance (Sec. 10.5). Truthprint instead hashes a *typed invariant contract* with a keyed MAC, a typed, keyed member of that family whose false-positive floor is the analyzed contract-collision probability (Sec. 10.4). Table 9 (appendix) positions it against representative systems.

3 Problem Definition

For a content specification x, secret key K, payload m, encoder E_K, language-ℓ realizer R_ℓ, transformation A, multilingual parser P, and detector D_K, the generated text is $y = R_\ell(E_K(x, m))$, an observed transform is $y' = A(y)$, and detection returns $(s, \hat{m}, c) = D_K(P(y'))$ (score, recovered payload or prefix, confidence report). The design targets four objectives: (i) *semantic fidelity* $I(P(y)) = I(x)$ for invariant extractor I; (ii) *natural-language quality* above deployment thresholds (not evaluated here); (iii) *robust detection* under transformations that preserve sufficient invariant and carrier evidence; and (iv) *controlled authenticated false positives*—under the stated PRF and independence assumptions an independently produced decoded payload rarely passes the keyed tag check, with the statistical fallback separately calibrated.

4 Threat Model

Benign transformations. The target is meaning-preserving edits (normalization, trans-

lation, synonym substitution, voice alternation, clause splitting, paraphrase); this paper evaluates translation, round-trip translation, and sentence-level paraphrase.

Attackers. A non-adaptive attacker applies off-the-shelf translation, paraphrasing, or summarization without knowing K. An adaptive attacker may know the schema and detector design, repeatedly rewrite text, splice documents, target known carrier families, or query an exposed detector.

Out of scope. Truthprint does not identify all AI-generated text. It verifies consistency with a participating keyed generator. Absence of a signal does not prove human authorship.

5 Truthprint Intermediate Representation

5.1 Three-Layer Structure

A Truthprint document contains three conceptual layers.

Layer A: Semantic invariants. Entities, predicates, semantic roles, polarity, negation scope, modality, quantities, temporal intervals, causal direction, conditions, exceptions, attribution, quotation status, and uncertainty.

Layer B: Controlled realization variables. Voice, event/state framing, clause packaging, discourse-connective family, proposition ordering, nominal/verbal realization, anaphoric choice, and optional explanatory apposition.

Layer C: Provenance metadata. Scheme and key identifiers, payload, nonce, ECC parameters, carrier assignments, and invariant digest.

5.2 Typed Mutability

Each field f is assigned $\text{Mutability}(f) \in \{\text{locked}, \text{conditional}, \text{realization-only}\}$: locked fields never change, conditional fields change only when context-specific validators approve, and realization-only fields may carry watermark symbols.

5.3 Example Representation

A simplified representation (locked predicate/roles/polarity/time with a realization space over voice and event framing) is given in Appendix J.

6 Formal Semantic Contract

A Truthprint graph is $G = (V_E, V_V, E_R, E_D, \Phi)$—entity nodes, event/state nodes, semantic-role edges, discourse/logical relations, and scoped features such as negation, modality, time, quantity, and attribution—partitioned as $G = I \cup C \cup U$ into the invariant set I, the controlled carrier set C, and unconstrained style U. A transformed structure G' is admissible only when $\text{InvariantEq}(I(G), I(G')) = 1$ and $Q(R_\ell(G')) \geq \tau_Q$ for a quality and contextual-consistency function Q. Hard invariants are validated symbolically; soft ones (emphasis, implicature) would need entailment checks.

7 In-Text Carrier Encoding

This in-text path fails under real MT (Sec. 10.1); the evaluated mechanism is the ledger tag over h_I (Sec. 10.4). The raw payload may carry a generator-family, deployment, coarse time-bucket, nonce, and policy identifier—never personal identifiers, prompts, or exact timestamps. The authenticated payload is $p = m \parallel \text{MAC}_K(m \parallel n \parallel h_I)$ for nonce n and canonical invariant digest h_I, error-corrected as $c = \text{ECC}(p)$ following prior robust multi-bit watermarking (Qu et al., 2025). The encoder discovers admissible carrier families per discourse unit (voice, clause packaging, discourse-connective family, nominal/verbal realization, proposition order), rejecting context-inappropriate ones such as a passive that obscures an important agent. For carrier j with alternatives $A_j = \{a_{j,0}, \ldots, a_{j,r-1}\}$, a keyed permutation $\pi_j = \text{PRP}_K(\text{CanonicalHash}(I, G, j, n))$ selects $a_j^\star = \pi_j(c_j \bmod |A_j|)$; because the map depends on the invariant graph and nonce, no global rule such as "passive means 1" exists. Each realized candidate is reparsed and checked for hard-invariant equality, bidirectional entailment, numerical and entity consistency, fluency, and expected carrier recoverability; if none is valid the carrier is recorded as an erasure rather than forcing degraded text. Encoding and detection pseudocode are in Appendix J.

8 In-Text Carrier Detection

Detection segments the document, identifies language, parses each segment into Truthprint,

resolves entities, normalizes predicates and roles, canonicalizes quantities and time, and extracts observable carrier alternatives. For an observed carrier a_j^{obs} it reconstructs the keyed permutation and derives $\hat{c}_j = \pi_j^{-1}(a_j^{obs})$, treating uncertain observations as erasures; an ECC decoder reconstructs candidate payloads, accepted only when their MAC verifies. When complete recovery is impossible, a calibrated fallback over n reliable binary carriers with k keyed matches uses

$$z = \frac{k - 0.5n}{\sqrt{0.25n}}, \tag{2}$$

whose production calibration must account for dependence among carrier observations.

9 Formal Analysis

We state what the design guarantees given sound primitives; what depends on semantic parsing is measured empirically (Sec. 10), and computational complexity is in Appendix I.

9.1 Carriers as Invariant-Equivalence Classes

Definition 1 (Invariant equivalence)**.** *Let $I(\cdot)$ be the invariant extractor. Two structures G, G' are invariant-equivalent, written $G \sim_I G'$, iff* $\mathrm{InvariantEq}(I(G), I(G')) = 1$. *For a discourse unit with invariants I, the admissible realization set is*

$$\mathcal{R}(I) = \{\, r : I(P(r)) = I \ \wedge\ Q(r) \geq \tau_Q \,\}. \tag{3}$$

A carrier alphabet A_j is a key-dependent partition of $\mathcal{R}(I)$ into classes; the encoder selects a class by symbol and the realizer emits any member. Thus the watermark ranges over the quotient $\mathcal{R}(I)/\!\sim_I$ rather than over tokens or embedding neighborhoods.

Assumption 1 (Validator soundness)**.** *The hard-invariant validator V is sound: if V accepts a candidate r for invariants I, then $I(P(r)) = I$. (It need not be complete; rejected admissible candidates only reduce capacity, not fidelity.)*

Remark 1 (Conditional protected-field fidelity; immediate from Assumption 1)**.** *Under Assumption 1, every committed output $y = R_\ell(E_K(x, m))$ satisfies $I(P(y)) = I(x)$.*

Proof sketch. Encoding commits a candidate only if V accepts it, and carriers with no accepted candidate are recorded as erasures rather than realized (Sec. 7). By soundness, acceptance implies invariant equality for that unit; taking the conjunction over committed units gives $I(P(y)) = I(x)$. □

Empirical validator soundness. We measure the load-bearing assumption as the validator's false-accept rate on single-field tampers. On the closed domain the typed contract flags every tamper across all six fields (property C1, 1.000, CI $[0.998, 1.000]$) while retaining attribution across benign paraphrase (0.990); a decode-only ablation accepts every tamper (C2, 0.000), confirming that invariant re-extraction, not the coding layer, supplies soundness. Under real NLLB translation a registered record altered in one core6 field fails to match the translated text at ≥ 0.996 (Sec. 10.4). Wide-coverage validator soundness remains open, so Remark 1 stays conditional.

9.2 Unforgeability and False Positives

Proposition 1 (Fresh-payload unforgeability; standard reduction)**.** *Model the tag as $t = \mathrm{MAC}_K(m \parallel n \parallel h_I)$ truncated to τ bits. Any PPT adversary that causes authenticated attribution for a tuple (m, n, h_I) not previously tagged by the generator yields an EUF-CMA forgery against this truncated MAC with the same success probability (up to simulation overhead). Thus its success is bounded by the EUF-CMA advantage of the instantiated truncated MAC; guessing loss is already part of that advantage and is not added a second time.*

Proof sketch. Attribution requires a recovered payload whose tag equals $\mathrm{MAC}_K(m \parallel n \parallel h_I)$ on a triple not produced by the generator. The reduction answers authorized generation requests with its tagging oracle and, when the detector accepts a fresh tuple, outputs that tuple and recovered tag as its forgery. Replay of a previously tagged tuple is excluded by freshness and must instead be handled by deployment policy. □

Proposition 2 (Authenticated acceptance under an ideal-PRF model)**.** *Assume the truncated MAC is a pseudorandom function on a fresh input and the decoded candidate tag is*

independent of its output. For a fixed (key, nonce) test, an independently produced candidate payload passes the τ-bit tag check with probability $2^{-\tau}$, up to the PRF distinguishing advantage. For Q tested key–nonce pairs, the union bound is $Q2^{-\tau}$ plus that advantage.

This bounds only a candidate that reaches tag verification: for q fresh candidates whose tags are independent of ideal-PRF outputs the union bound is $q2^{-\tau}$ (2.3×10^{-10} at τ=32), *not* an end-to-end detector FPR—parsing, candidate search, replay, carrier dependence, and threshold calibration all affect system-level error, and replayed outputs require a nonce registry. The 0/20000 observation of Sec. A is consistent with the bound (approx. 95% binomial upper bound 1.5×10^{-4}). When full recovery fails, the fallback z of Eq. (2) satisfies $\Pr[z \geq t \mid H_0] \leq e^{-t^2/2}$ under independent binary carriers, calibrated in deployment.

Carrier capacity and the minimum reliable text length for the in-text channel are analyzed in Appendix H.

10 Experiments

Setup. All data are *synthetic template facts* in one closed domain (software incidents: four agents, four patients, one predicate—*fix*, with *break* used only as a tamper value—and two to four closed values each for polarity, quantity, time direction, modality, attribution, and causation, so core6 has 192 possible values); the main corpus is 50 documents ×16 sentences, translated by NLLB-200-distilled-600M into Korean, Hindi, Chinese, Arabic, and German and round-trip into English (4,800 outputs). The *Stage-1* frontend parses surface carriers; the *Stage-2* frontend is a hand-built six-language lexicon that extracts the nine invariant fields. A *contract* is the field subset that is authenticated: core6 (polarity, quantity, time, attribution, predicate, agent) or full9 (all nine). Because the extractor and the corpus share one closed vocabulary, these results establish feasibility, not open-domain accuracy. Runs are seeded; rates carry Wilson 95% CIs; implementation, protocol, and hypotheses are in the appendix.

10.1 Surface carriers: a negative result under real MT

On the committed run, the Stage-1 *surface* frontend reads all 800 English source sentences but only 1 of the 4,800 NLLB translations (parse coverage 0.000, CI $[0.000, 0.001]$; the one success is a round trip). Hiding a payload in realization freedom therefore does not survive real translation; the bottleneck is wide-coverage parsing, not the coding or cryptographic core (Appendix A).

10.2 The meaning layer survives real translation (Stage-2)

A closed-domain lexicon extractor over six languages (English, Korean, Hindi, Chinese, Arabic, German) recovers the nine invariant fields from 4,800 real NLLB translations (800 source sentences over 50 documents × six conditions). Aggregate recovery (Wilson 95% CI) is 0.83–0.99: the small closed categoricals (polarity, quantity, attribution, time) survive at ≥ 0.97, while the open entity and relational fields are lower (0.83–0.92; patient 0.826 weakest). Strict all-nine-field recovery is highest for Hindi/German (0.69/0.66) and lowest for Arabic (0.26; Arabic loses the patient *cache miss*, Appendix G.3). The meaning layer thus survives MT where the surface carriers of Sec. 10.1 do not (`scripts/eval_multilingual.py`).

10.3 A wide-coverage neural frontend on open-domain text

A feasibility pilot (16 base sentences, 6 with open wording, in six conditions; 96 records) ran an instruction LLM as a drop-in extractor: on *closed* templates the hand-built lexicon leads (0.970 vs. 0.726 field recovery), but on the open wording the lexicon degrades below the neural frontend (0.481 vs. 0.614), strongest on typed relational fields (predicate 0.896, causation 0.948). Neural extraction is not yet reliable: it recovers all nine fields of a sentence in none of the 96 records (lexicon 0.479). The harness enforces *separate* translation and extraction models (removing the self-consistency confound).

10.4 Meaning-based provenance authentication under real translation

The producer registers a keyed τ-bit tag (τ=32) over the canonical digest of a declared *invariant contract* (a field subset) in the out-of-band ledger; the verifier re-extracts the invariants (Stage-2) and authenticates iff the tag reproduces. The key matters only if the ledger is exposed: core6 has 192 values, so an unkeyed digest could be inverted by enumeration and new entries minted by anyone, whereas a tag needs the key (Proposition 1); with a private, trusted ledger a plain digest decides identically (not evaluated). On the 4,800 translations a compact contract (polarity, quantity, time, attribution, predicate, agent) gives document-level attribution—a document attributed when $\geq$ half its sentences authenticate—of 1.000 (CI $[0.929, 1.000]$) for German and Hindi, 0.980 round-trip, 0.940 Korean/Chinese, and 0.900 $[0.786, 0.957]$ Arabic over 50 documents (unchanged after correcting a generator plural bug and re-translating; App. G.3); sentence-level TPR 0.619–0.901, tamper rejection $\geq$ 0.996, false positives 0.000–0.007. Tampers alter one *contract* field of the registered record (no edited text is re-translated); edits outside the contract (core6: patient, modality, causation) pass by design. Widening to nine fields raises meaning coverage but lowers TPR (AR 0.256) and collapses attribution for the hardest conditions (AR 0.02, ZH 0.08, RT 0.28) while German/Hindi hold at 0.94—an explicit fidelity/robustness trade-off. Two low-entropy generations sharing a *contract* collide before the tag is tested, so the false-positive floor is the *contract-collision* probability, not $2^{-\tau}$. For the six-field contract this floor is 5.2×10^{-3} ($\approx 1/192$), matching the measured 0.000–0.007 and $\sim 2 \times 10^{7}\times$ the 2^{-32} bound; nine fields lower it to $\sim 1.8 \times 10^{-4}$, and the $2^{-\tau}$ bound binds only once the entity-vocabulary entropy exceeds τ bits (Appendix A). These rates are *id-addressed*: the verifier is told which registered sentence to test. Without an id (translation strips any in-text identifier) the verifier must search the ledger, and sentence-level search is not specific—a core6 sentence collides with some wrong entry of the 800-sentence ledger in 0.66–0.89 of cases. The $\geq$50% document rule restores specificity: searching all 50 registered documents with no id, core6 attributes translated documents to their source at 0.90–1.00 with 0/2,450 false document attributions per condition (95% upper bound 1.6×10^{-3} per document pair). Per-pair risk multiplies with ledger size, which we measure by adding up to 10^5 distractor documents from the corpus generator (only queries need translation): core6 makes no false attribution up to 1,050 documents, but 1–6 per condition at 10,050 and 3–61 at 100,050 ($\approx 5 \times 10^{-6}$ per pair, near the independence prediction of 1.1×10^{-5}), whereas full9 makes none up to 100,050 documents (5×10^6 pairs per condition). At 10^5 documents no contract is both robust and specific: core6 falsely attributes 6–32% of query documents, full9 attributes only 0.02–0.94. The NLLB run is in the supplementary material (`paper/results/realmt_b2/`); every number here reproduces with no GPU (`scripts/eval_provenance.py`, `eval_ledger_search.py`, `eval_ledger_scale.py`).

10.5 Baseline comparison: token, semantic, and robust-hashing

A KGW token watermark with γ=0.25, δ=2.0 (Kirchenbauer et al., 2023) on GPT-2 generations, translated by the same NLLB system, loses document-level detection from 0.960 clean to 0.000–0.380 across six conditions (Appendix Table 12; cf. (He et al., 2024)). On the *same* text (30 documents reworded by Qwen2.5-0.5B with KGW, SynthID-Text, or no watermark; Appendix Table 13) both token watermarks fall to 0.000–0.133 after KO/AR/DE translation, while Truthprint attributes 0.633–0.967 with no false attribution—with caveats: Truthprint alone resamples (up to four times) until its record reads back, 12.9% of sentences still fail, and SynthID-Text is scored by its mean g-value, not a trained detector (clean TPR only 0.267). The controlled comparison scores every method at a *matched* 1% FPR on the same 800 NLLB source sentences and their translations (300 bootstrapped 16-sentence documents per condition; static multilingual embedding `model2vec`; Table 1). Its document score is the fraction of id-addressed

Table 1: Matched-FPR comparison on the 800-sentence NLLB corpus: fraction of translated 16-sentence documents identified at a common 1% FPR (300 bootstrapped documents per condition; they reuse sentences, so are not independent). RT = round trip via Korean; SemStamp = our re-implementation; * = paraphrase-tuned encoder, else static. Only the typed contract also rejects single-field tampers (in text).

Method	RT	KO	HI	ZH	AR	DE
Typed	1.000	1.000	1.000	1.000	1.000	1.000
SemStamp	0.300	0.020	0.057	0.053	0.023	0.020
SimHash	0.973	0.293	0.240	0.773	0.050	0.450
SemStamp*	0.213	0.110	0.090	0.060	0.030	0.053
SimHash*	0.997	0.990	1.000	1.000	1.000	1.000
MinHash	0.737	0.000	0.000	0.000	0.000	0.000
SHA-256	0.020	0.000	0.000	0.000	0.000	0.000

sentences that verify, thresholded at the calibrated FPR (1–2 of 16 for the typed digest); Sec. 10.4 and Table 13 instead require half the sentences, hence their lower attribution. The typed meaning-digest identifies every translated document in all six conditions (1.000); a SemStamp-style LSH detector (Hou et al., 2024)—our re-implementation, granted ideal generation—reaches 0.020–0.300, embedding SimHash (Charikar, 2002) 0.050–0.973, and MinHash (Broder, 1997) and SHA-256 nothing beyond the round trip. With a paraphrase-tuned multilingual encoder (Reimers and Gurevych, 2020) SemStamp-style stays at 0.030–0.213, but SimHash survives translation (0.990–1.000): a strong encoder closes the *survival* gap. Near-duplicate templates handicap few-bit regions (on a separate 16-sentence LLM-MT set SemStamp-style reaches 0.377–1.000; Appendix Table 11). The difference that remains is tampering: a single-field meaning-altering edit is rejected by the typed contract (1.000, CI $[0.981, 1.000]$) but by *neither* similarity hash (0.000; 0.005 with the paraphrase encoder)—a near-duplicate fingerprint cannot see which typed field changed—while the exact hash that rejects every tamper also rejects every legitimate translation (Table 1).

10.6 Ledger retrieval and NLI baselines

Because the scheme already uses an out-of-band ledger, the natural control is the established retrieval defense (Krishna et al., 2023): store the source and *retrieve* it after transformation (bag-of-words cosine nearest neighbor), optionally with an NLI check. Under the same ledger and id-addressed verification, at benign acceptance of 0.94/0.92 (typed digest 1.000), plain retrieval and a directional-coverage NLI proxy reject essentially no single-field tamper (0.047/0.061; 0.186/0.195 with a real multilingual encoder): a one-field edit stays nearest its own source. Only a *perfect*, unrealizable oracle NLI matches the typed contract's 1.000 rejection, and needs a semantic model the typed digest does not (Appendix Table 14).

10.7 Paraphrase and adaptive-attack robustness

Rule-based, in-lexicon paraphrase (voice, time reflow, verb/time synonyms, hedge) authenticates at TPR 1.000 and document attribution 1.000, and a single-field altering edit is rejected at 1.000. A schema-aware adaptive attacker that swaps the predicate verb and time expression for *out-of-lexicon* synonyms achieves ValidRemoval (Eq. (1)) 1.000 under the *closed* lexicon—a negative result we localize to exactly the abstained fields (predicate and time in all cases; attribution when present), with polarity and agent recovered in 100% of attacks. The vulnerability is the fixed lexicon, not the meaning-digest: an *extended-coverage* frontend drives it $1.000 \to 0.000$ at unchanged benign and tamper rejection, but a *held-out* out-of-inventory attack still evades any fixed lexicon (1.000). Two small runs with real instruction LLMs as open-vocabulary parsers cut the held-out ValidRemoval to 0.050 (GPT-4o-mini, N=40, CI $[0.014, 0.165]$) and 0.100 (Llama-3.3-70B, N=10, CI $[0.018, 0.404]$) at benign and tamper rejection 1.000, though Llama still loses 0.200 of the *seen* attacks that the extended lexicon stops, suggesting open-vocabulary parsing as the general defense (Appendix Tables 10, 15).

11 Conclusion

Truthprint separates immutable meaning from controlled realization with a machine-checkable typed contract; what survives translation is a keyed tag over that contract, re-verified out of band—a complement to token-level watermarks.

Limitations

Multilingual semantic parsing remains imperfect. Not every proposition admits multiple safe realization alternatives. Highly constrained text has low carrier capacity. Some apparently stylistic alternations change emphasis or information structure. Translation can omit modality, evidentiality, subjects, or discourse relations. Strong summarization can destroy both content and watermark evidence. Finally, candidate generation and semantic validation add overhead compared with direct token-level methods: on the same-text run the validation loop needs 1.79 draws per sentence on average and still finds no valid realization for 12.9% of sentences (Appendix Table 13). A single tag over the contract digest says *whether* a protected field changed, not *which*; naming the field would need per-field tags or a stored record, which we did not evaluate and which increases exposure. Verification also requires the generator's ledger and, because the tag is a MAC, its secret key: a third party cannot check provenance on its own, and public verification would need a digital signature over the digest (with longer tags), which we have not evaluated. The closed-domain adaptive paraphrase of Section 10.7 shows the sharpest limitation directly: a schema-aware attacker who swaps out-of-lexicon content synonyms fully removes the meaning-digest under the closed lexicon while preserving meaning. An extended-coverage frontend defends that specific attack (ValidRemoval $1.000 \rightarrow 0.000$ without weakening tamper rejection); a held-out out-of-inventory attack still evades any fixed lexicon, but two real open-vocabulary neural runs (GPT-4o-mini N=40, Llama-3.3-70B N=10; Table 10) resist even the held-out attack (0.050/0.100), suggesting that wide-coverage *open-vocabulary* parsing—not a longer lexicon—is the general defense. Those runs are closed-domain and modest in scale; the current artifact still does not evaluate real generators, human judgments, all official baselines, replay defenses, or open-vocabulary adaptive attacks *at larger scale and across domains.* All evaluation data are synthetic template facts in one closed domain, read by a hand-built extractor that shares the corpus vocabulary; the matched-FPR comparison uses our own SemStamp-style re-implementation on template sentences that handicap few-bit embedding regions (with static and paraphrase-tuned encoders); the same-text token-watermark control uses one small generator (30 documents), and Truthprint's lower KO/AR attribution there (0.633/0.667) shows that freer LLM wording strains the extractor; and id-free ledger attribution was measured against ledgers of up to 10^5 documents drawn from the same low-entropy template generator, not real text. Consequently it supports feasibility and closed-domain robustness claims, not comparative robustness or deployment readiness at wide coverage.

Ethical Considerations

A provenance detector can be misused to accuse authors, censor content, or make automated disciplinary decisions. Truthprint therefore treats a positive result as evidence about a participating keyed generation system, not proof of AI authorship or of misconduct. Deployments should expose the tested scheme and key identifier, use independently calibrated error rates, preserve an appeal path, and avoid adverse decisions based on the detector alone. Absence of a mark is likewise not evidence of human authorship.

The design has dual-use risks: public carrier families can aid removal attacks, while secret or overly detailed payloads can enable tracking. The ledger retains a typed record of every registered generation, and a deterministic tag reveals when two outputs share a contract, so it must be access-controlled and minimized. Payloads should exclude personal identifiers, prompts, and precise timestamps; keys and nonce registries require access control, rotation, and audit. Evaluation text should be licensed or consented, human annotation should be compensated and documented, and language-specific error disparities should be reported rather than hidden in an aggregate score. The present artifact uses synthetic template facts, no personal data, and no human subjects; the real-MT pilot (Secs. 10.1–10.4) translates those synthetic facts with a public open-weights model (NLLB-200-distilled-600M) and involves no personal data or human subjects. Any scaled evaluation that adds human invari-

ant or factual-equivalence annotation must use licensed or consented text and compensated, documented annotation.

References

Andrei Z. Broder. 1997. On the resemblance and containment of documents. In *Proc. Compression and Complexity of Sequences*, pages 21–29.

Moses S. Charikar. 2002. Similarity estimation techniques from rounding algorithms. In *Proc. ACM STOC*, pages 380–388.

Yize Cheng, Vinu Sankar Sadasivan, Mehrdad Saberi, Shoumik Saha, and Soheil Feizi. 2025. Adversarial paraphrasing: A universal attack for humanizing AI-generated text. In *Proc. NeurIPS*.

Miranda Christ, Sam Gunn, and Or Zamir. 2024. Undetectable watermarks for language models. In *Proc. COLT*.

Sumanth Dathathri, Abigail See, Sumedh Ghaisas, Po-Sen Huang, Rob McAdam, Johannes Welbl, Vandana Bachani, Alex Kaskasoli, Robert Stanforth, Tatiana Matejovicova, Jamie Hayes, Nidhi Vyas, Majd Al Merey, Jonah Brown-Cohen, Rudy Bunel, Borja Balle, Taylan Cemgil, Zahra Ahmed, Kitty Stacpoole, and 5 others. 2024. Scalable watermarking for identifying large language model outputs. *Nature*, 634(8035):818–823.

European Parliament and Council. 2024. Regulation (EU) 2024/1689 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act). Official Journal of the European Union.

Pierre Fernandez, Antoine Chaffin, Karim Tit, Vivien Chappelier, and Teddy Furon. 2023. Three bricks to consolidate watermarks for large language models. In *Proc. IEEE WIFS*.

Thibaud Gloaguen, Nikola Jovanović, Robin Staab, and Martin Vechev. 2025. Black-box detection of language model watermarks. In *Proc. ICLR*.

Google. 2026. SynthID: Tools for watermarking and detecting LLM-generated text. Responsible Generative AI Toolkit, Google AI for Developers.

Google DeepMind. 2024. SynthID-Text reference implementation. GitHub repository.

Xia Han, Qi Li, Jianbing Ni, and Mohammad Zulkernine. 2025. Robustness assessment and enhancement of text watermarking for Google's SynthID. In *Proc. IEEE TrustCom*.

Zhiwei He, Binglin Zhou, Hongkun Hao, Aiwei Liu, Xing Wang, Zhaopeng Tu, Zhuosheng Zhang, and Rui Wang. 2024. Can watermarks survive translation? on the cross-lingual consistency of text watermark for large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 4115–4129.

Abe Bohan Hou, Jingyu Zhang, Tianxing He, Yichen Wang, Yung-Sung Chuang, Hongwei Wang, Lingfeng Shen, Benjamin Van Durme, Daniel Khashabi, and Yulia Tsvetkov. 2024. SemStamp: A semantic watermark with paraphrastic robustness for text generation. In *Proc. NAACL*.

John Kirchenbauer, Jonas Geiping, Yuxin Wen, Jonathan Katz, Ian Miers, and Tom Goldstein. 2023. A watermark for large language models. In *Proc. ICML*.

Kalpesh Krishna, Yixiao Song, Marzena Karpinska, John Wieting, and Mohit Iyyer. 2023. Paraphrasing evades detectors of AI-generated text, but retrieval is an effective defense. In *Advances in Neural Information Processing Systems (NeurIPS)*.

Rohith Kuditipudi, John Thickstun, Tatsunori Hashimoto, and Percy Liang. 2024. Robust distortion-free watermarks for language models. *Transactions on Machine Learning Research*.

Aiwei Liu, Leyi Pan, Xuming Hu, Shu'ang Li, Lijie Wen, Irwin King, and Philip S. Yu. 2024a. An unforgeable publicly verifiable watermark for large language models. In *Proc. ICLR*.

Aiwei Liu, Leyi Pan, Xuming Hu, Shiao Meng, and Lijie Wen. 2024b. A semantic invariant robust watermark for large language models. In *Proc. ICLR*.

Aiwei Liu, Leyi Pan, Yijian Lu, Jingjing Li, Xuming Hu, Xi Zhang, Lijie Wen, Irwin King, Hui Xiong, and Philip S. Yu. 2024c. A survey of text watermarking in the era of large language models. *ACM Computing Surveys*, 57(2):1–36.

NLLB Team, Marta R. Costa-jussà, James Cross, Onur Çelebi, Maha Elbayad, Kenneth Heafield, Kevin Heffernan, Elahe Kalbassi, Janice Lam, Daniel Licht, Jean Maillard, Anna Sun, Skyler Wang, Guillaume Wenzek, Al Youngblood, Bapi Akula, Loic Barrault, Gabriel Mejia Gonzalez, Prangthip Hansanti, and 20 others. 2024. Scaling neural machine translation to 200 languages. *Nature*, 630(8018):841–846.

Romina Omidi, Yun Dong, and Binghui Wang. 2026. On Google's SynthID-Text LLM watermarking system: Theoretical analysis and empirical validation. *Preprint*, arXiv:2603.03410.

Leyi Pan, Aiwei Liu, Zhiwei He, Zitian Gao, Xuandong Zhao, Yijian Lu, Binglin Zhou, Shuliang Liu, Xuming Hu, Lijie Wen, Irwin King, and Philip S. Yu. 2024. MarkLLM: An open-source toolkit for LLM watermarking. In *Proc. EMNLP (System Demonstrations)*.

Wenjie Qu, Wengrui Zheng, Tianyang Tao, Dong Yin, Yanze Jiang, Zhihua Tian, Wei Zou, Jinyuan Jia, and Jiaheng Zhang. 2025. Provably robust multi-bit watermarking for AI-generated text. In *Proc. USENIX Security*.

Nils Reimers and Iryna Gurevych. 2020. Making monolingual sentence embeddings multilingual using knowledge distillation. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 4512–4525.

Jie Ren, Han Xu, Yiding Liu, Yingqian Cui, Shuaiqiang Wang, Dawei Yin, and Jiliang Tang. 2024. A robust semantics-based watermark for large language model against paraphrasing. In *Findings of the Association for Computational Linguistics: NAACL 2024*, pages 613–625.

Tom Sander, Pierre Fernandez, Alain Durmus, Matthijs Douze, and Teddy Furon. 2024. Watermarking makes language models radioactive. In *Proc. NeurIPS*.

Jonas Schäfer, Cezary Pilaszewicz, and Gerhard Wunder. 2026. Robust text watermarking for large language models via dual semantic embeddings. In *Findings of the Association for Computational Linguistics: EMNLP 2026*.

Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Nicolas Papernot, Ross Anderson, and Yarin Gal. 2024. AI models collapse when trained on recursively generated data. *Nature*, 631(8022):755–759.

Ziping Ye, Gourab Dey, Christos Christodoulopoulos, Charith Peris, Anil Ramakrishna, Weitong Ruan, Aram Galstyan, Kai-Wei Chang, Rahul Gupta, and Ninareh Mehrabi. 2026. SWAN: Semantic watermarking with Abstract Meaning Representation. In *Proc. ACL*.

Xin Yi, Yue Li, Shunfan Zheng, Linlin Wang, Xiaoling Wang, and Liang He. 2025. Unified attacks to large language model watermarks: Spoofing and scrubbing in unauthorized knowledge distillation. *Knowledge-Based Systems*, 329:114295.

Xuandong Zhao, Prabhanjan Ananth, Lei Li, and Yu-Xiang Wang. 2024. Provable robust watermarking for AI-generated text. In *Proc. ICLR*.

A Reference Implementation and Verification

To make the guarantees above concrete and reproducible, we provide a self-contained reference implementation (`reference/truthprint_poc.py`, standard library only) of the language-independent core: the canonical invariant digest h_I, the authenticated payload $p = m \parallel \mathrm{MAC}_K(\cdot)$, the invariant-bound keyed carrier map, and a systematic GF(2) code with erasure decoding. This deliberately excludes the semantic parser so that the cryptographic and coding properties can be validated in isolation from model-dependent NLP. It exercises four properties and reports the following verified behavior (message $\mu = 16$ bits, tag $\tau = 32$ bits, code $[N, k] = [96, 48]$):

- **P1 (round-trip recovery).** An invariant-preserving transform that re-realizes carriers and erases 28/96 sites still yields exact payload recovery and a verifying MAC.
- **P2 (invariant binding).** Flipping a locked field (polarity positive→negative) changes h_I, so the MAC fails and the document is *not* attributed, illustrating Remark 1's contract.
- **P3 (cryptographic false positives).** Over 20,000 independent documents, 0 were attributed. This is consistent with Proposition 2; statistically, the observation alone only gives an approximate 95% binomial upper bound of 1.5×10^{-4}.
- **P4 (no global rule).** For a fixed symbol, $43/96 \approx N/2$ realized options differ between two documents whose invariants differ, confirming the map is bound to (K, h_I, n).

Table 2 reports decode success versus carrier erasure rate for the implemented $[96, 48]$ code and seeded random-erasure protocol: full recovery (≥ 0.998) through $\epsilon = 0.42$, followed by 0.982 at $\epsilon = 0.44$, 0.774 at $\epsilon = 0.48$, 0.290 at $\epsilon = 0.50$, and zero observed recovery at $\epsilon \geq 0.52$. The location and shape are finite-code measurements; only the necessary survivor-count boundary follows from Proposition 3. These values characterize the coding layer, not detection on model-generated text.

A.1 Closed-Domain Linguistic Round-Trip

The core PoC operates on abstract carrier bits. To validate the linguistic layer end to end, a second self-contained script (`reference/truthprint_linguistic_demo.py`) realizes an authenticated payload across templated factual sentences using two invariant-preserving carriers—voice (active/passive) and time-phrase position—then rewrites, re-parses, and

Table 2: Reference-implementation decode success vs. carrier erasure rate ($[96, 48]$ GF(2) code, 500 trials per point, seeded). Fine resolution near the cliff reveals a gradual transition from full recovery at $\epsilon = 0.42$ to collapse at $\epsilon \geq 0.52$, as predicted by Proposition 3.

Carrier erasure ϵ	Recovery rate
0.10	1.000
0.20	1.000
0.30	1.000
0.40	1.000
0.42	0.998
0.44	0.982
0.45	0.968
0.46	0.950
0.48	0.774
0.50	0.290
0.52	0.000
0.55	0.000
0.60	0.000

detects on actual sentence strings (the running example of Section E). On a 32-sentence document (64 carriers, code $[64, 32]$, $\tau = 20$) it verifies:

- **L1 (invariant fidelity).** Every benign transform leaves the parsed agent, predicate, patient, and polarity unchanged.
- **L2 (round-trip recovery).** With 19/64 carriers erased by realization-canonicalizing rewrites, the exact payload is recovered and the MAC verifies.
- **L3 (tamper detection).** Negating a single fact changes the document invariant digest, so attribution correctly fails.

This is a Stage-1 closed-domain demonstration, not a general multilingual parser; it establishes that the contract, keyed carriers, and erasure-aware decoder compose correctly once realized as text, and it isolates the remaining open problem as robust wide-coverage semantic parsing (Section B).

A.2 Field-Level Challenge Set and Component Ablation

To test—rather than assume—that a *typed* invariant contract catches what an embedding *similarity* judge misses, we extend the closed-domain frontend to six independently protected fields (polarity, quantity, temporal direction, modality, attribution, causation) and generate minimal-pair probes (`reference` is packaged as `truthprint.challenge`; regenerate with `scripts/eval_challenge.py`). For each field we compare a meaning-preserving realization change against a single-field meaning-altering edit, and judge each with three detectors that differ only in the admissibility rule: the full typed+MAC contract, a decode-only ablation that drops the h_I-bound MAC, and an embedding-similarity gate (bag-of-words cosine) *calibrated* to reject at most 5% of legitimate paraphrases. On 300 documents of 24 facts ($[48, 24]$ code), with 95% confidence intervals (Wilson for rates, bootstrap for cosines) at a fixed seed, we measure:

- **C1 (typed tamper detection).** The typed+MAC contract detects every single-field tamper across all six fields (aggregate 1.000, CI $[0.998, 1.000]$) while retaining attribution across benign paraphrases (0.990, CI $[0.971, 0.997]$).
- **C2 (MAC is load-bearing).** The decode-only ablation detects *no* tamper (aggregate 0.000, CI $[0.000, 0.002]$): the carriers still recover, so the h_I-bound MAC—not the coding layer—is what resists a meaning-altering edit.
- **C3 (similarity cannot separate meaning).** The calibrated embedding gate catches almost no tamper (aggregate 0.012, CI $[0.008, 0.018]$), because a single-field edit is lexically smaller than a legitimate voice/time reflow: in 48.6% of cases (CI $[0.463, 0.509]$) the tamper is at least as cosine-similar to the source as the paraphrase is, so *no* similarity threshold can separate them.

This isolates the interaction the design relies on—typed invariants supply the decision an embedding metric cannot, and the MAC binding converts that decision into tamper resistance—on real strings. It remains a closed-domain result; the wide-coverage multilingual analogue requires the neural frontend and human annotations specified in Section B and the accompanying data templates.

Table 3: Evaluation matrix: each research question mapped to its decisive comparison, manipulated factor, and primary outcome measure.

RQ	Decisive comparison	Primary measure
RQ1	Truthprint vs. unconstrained paraphrase watermark	Invariant error rate (IER), contradiction rate
RQ2	Truthprint vs. token/embedding baselines, clean text	TPR at $\{1, 0.1, 0.01\}$% FPR; ROC-AUC
RQ3	Pre- vs. post-translation (1-hop, round-trip)	Payload BER, partial-recovery rate
RQ4	Sentence- vs. paragraph-level paraphrase	TPR@1%FPR vs. edit distance
RQ5	Watermarked vs. unwatermarked generation	Quality delta, added latency
RQ6	Recovery vs. document length	Authenticated bits recovered
RQ7	Adaptive attacker with vs. without invariant constraint	ValidRemoval rate (Eq. (1))

B Experimental Methodology

B.1 Research Questions

The original research questions, with their status in this paper, are:

- RQ1: Does Truthprint preserve immutable meaning better than unconstrained paraphrase-based watermarking?
- RQ2: What true-positive rate is achieved at fixed false-positive rates on clean text?
- RQ3: How much signal survives single-hop and round-trip translation?
- RQ4: How robust is Truthprint to sentence- and paragraph-level paraphrasing?
- RQ5: What quality and latency costs are introduced?
- RQ6: How many authenticated payload bits can be recovered as document length increases?
- RQ7: Can an adaptive attacker remove the watermark while preserving all locked invariants?

Table 3 fixes, for each question, the single manipulated factor and the outcome that would confirm or falsify the corresponding hypothesis, so that results cannot be reported selectively. *Status.* RQ3 is answered for the ledger path on real NLLB output (Secs. 10.2–10.4) and negatively for in-text carriers (Sec. 10.1), not via payload BER. RQ2 is answered by the matched-FPR and same-text comparisons (Tables 1, 13) after translation, not at 0.1/0.01% FPR. RQ4 and RQ7 (Eq. 1) are measured on the closed domain (Section 10.7). RQ1, RQ5, and RQ6 were *not* measured: there is no quality, latency, or human evaluation and no recovery-vs-length study on real text.

B.2 Diagnostic Simulation (Not an End-to-End Benchmark)

A controlled Stage-1 simulator checks the evaluation code paths and illustrates, under explicitly stipulated channel parameters, how signal placement (token vs. meaning layer) reacts to surface replacement and carrier erasure. Because the "baselines" are method-inspired proxies run from a shared harness rather than the published systems, the simulator *must not be used to rank methods* and is not evidence of real-translation performance. To keep this diagnostic from being mistaken for a benchmark, its full protocol and outputs (Tables 4 and 5) are deferred to Appendix F; the measured results this paper relies on are the real-MT findings of Sections 10.1–10.4 and the coding-core verification of Section A.

C Falsifiable Hypotheses

These hypotheses guided the design. H1 is *not supported* for in-text carriers, which fail under real MT (Sec. 10.1); its ledger counterpart is supported only at small, closed-domain scale (Tables 1, 13). H2–H6 are untested on wide-coverage text:

- H1: Truthprint will outperform token-level watermarking under translation because the carrier is reconstructed from language-independent semantics.
- H2: Explicit invariant validation will reduce errors in negation, quantity, time, causality, and attribution.
- H3: Truthprint may underperform token methods on very short text because seman-

tic carrier capacity accumulates across sentences.

- H4: ECC and erasure-aware decoding will improve recovery under parsing errors and partial deletion.
- H5: Strong summarization will remain difficult because it removes propositions and their associated carriers.
- H6: Combining structured carriers with embedding-based fallback evidence will outperform either alone across diverse languages.

D Security and Deployment Considerations

Carrier mappings should depend on the secret key, invariant digest, document nonce, and carrier identity. This prevents simple global carrier rules and reduces forgery risk. Exact detector scores should be access controlled or coarsened to reduce oracle attacks. The detector should localize evidence by paragraph and avoid attributing a mixed document to one source.

The reference API supplies the nonce to the detector out of band; natural-language carrier recovery alone does not discover it. A deployment must therefore define an authenticated nonce-resolution mechanism, such as a bounded time-bucket search plus deployment identifier or a provenance-ledger lookup. Every tested key–nonce pair consumes false-positive budget and must be included in calibration. A ledger should reject duplicate nonces and distinguish exact replay from fresh attribution. Splicing requires finer-grained binding (for example, paragraph identifiers and a document Merkle root); without it, a valid paragraph proves only that the participating generator produced that paragraph, not the provenance of the surrounding document.

Watermarks should support transparency rather than automated punishment. Every detection report should include the tested key and scheme version, text length, number of reliable carriers, calibrated false-positive rate, language and attack assumptions, confidence, and MAC verification status. User identifiers, raw prompts, and exact timestamps should not be embedded.

D.1 Deployment Path and Graceful Degradation

Truthprint is designed to slot into a serving stack as a semantic realization pass rather than as a modification of the token sampler. Two integration modes are practical. In the *inline* mode, the generator emits a draft, the encoder re-realizes admissible discourse units to carry payload symbols, and the validator gates each candidate; although the coding and cryptographic layers are independent of model size (Sec. 9), the unmeasured generation and validation overhead may be substantial. In the *asynchronous* mode, watermarking runs as a post-generation service, which decouples watermark latency from the user-facing generation latency at the cost of a short delay before provenance is finalized.

Because capacity is unevenly distributed across text (Eq. 5), a deployable system must degrade gracefully. Low-capacity spans—code, formulas, legal clauses, and verbatim quotations—admit few invariant-preserving realizations and are therefore poor carriers. For such spans Truthprint records erasures and can optionally fall back to a token-level watermark such as SynthID-Text, yielding a layered scheme intended to combine semantic-layer transformation tolerance with token-layer short-span coverage; both properties require end-to-end evaluation. A deployed detector should report evidence per paragraph, report the number of reliable carriers alongside the calibrated false-positive rate, and abstain rather than guessing when the recovered capacity falls below the minimum reliable length. This makes the system's competence explicit at detection time, which is the behavior an auditor or platform integrator requires.

E Worked Example

Consider the fact with locked invariants FIX(agent=developer, patient=server_error), polarity positive, time *previous_day*. Two binary carriers are available: VOICE $\in$ {active, passive} and TIME-POS $\in$ {front, end}. Suppose the keyed map assigns the code symbols so that this unit must realize as passive with the time phrase at the front:

> *"On the previous day, the server error was fixed by the developer."*

A benign transformation—say an English→Korean translation—typically preserves the invariants but may not preserve every carrier. If translation canonicalizes voice (Korean commonly renders this actively), the VOICE carrier is destroyed; the detector, on re-parsing, flags it as *low confidence* and records an *erasure* rather than reading a wrong bit. The TIME-POS carrier and the carriers of neighboring sentences survive. Because the payload is spread across sentences and protected by an error-correcting code, the authenticated payload is still recovered once the surviving carriers exceed the capacity floor of Eq. (5). In contrast, a rewrite that inserts a negation ("was *not* fixed") alters a locked field, changes the invariant digest h_I, and correctly fails MAC verification. This example is executed end-to-end on sentence strings in the reference implementation (Section A, properties L1–L3).

F Diagnostic Simulation: Full Protocol and Outputs

This appendix gives the full protocol and outputs of the controlled Stage-1 diagnostic summarized in the main text. It validates the simulator implementation only; it is not a neural or translation benchmark and must not be used to rank the named published systems.

Tables 4 and 5 report simulator outputs for Truthprint *and* four baselines on a shared closed-domain Stage-1 testbed (500 documents, 32 sentences each). To keep the comparison apples-to-apples in the absence of a wide-coverage neural frontend, every method is evaluated under one identical meaning-preserving channel and one identical protocol; the baselines are method-inspired proxies for selected aspects of each published method's signal placement, run from the same harness (`scripts/eval_baselines.py`), and abstract the neural frontend (LM sampler, sentence encoder, AMR parser) in exactly the way the Truthprint Stage-1 numbers do. This diagnostic checks code paths and illustrates hypotheses under explicitly chosen channel parameters. It is not evidence that any named method attains the reported performance on real translation, and it must not be used to rank the methods. End-to-end evaluation with real generators, translation systems, parsers, official baseline implementations, held-out calibration data, and confidence intervals remains future work.

Shared channel. A single meaning-preserving transform is characterised by two physical quantities that all methods see identically: the fraction τ_{tok} of surface tokens the transform replaces (cross-script translation replaces almost all tokens) and the fraction ϵ_{inv} of meaning-layer carriers lost to parse normalisation (the quantity Truthprint's own analysis uses). Each method's survival is then *derived* from these two numbers through its own signal dependence, so the relative ordering is forced by where the signal lives, not tuned per method: a KGW-style token-identity green-list proxy survives at $(1-\tau_{\text{tok}})^2$; an edit-aligned token watermark (DEW) at $(1-\tau_{\text{tok}})$ up to an edit budget then collapses; a sentence-embedding region (SemStamp) flips near the LSH boundary at a rate $\propto \tau_{\text{tok}}$ but is meaning-dominant; and meaning-layer signals (SWAN, Truthprint) are unaffected by τ_{tok}, losing only the ϵ_{inv} parse fraction.

Protocol. For each method a decision threshold is calibrated at 1% FPR on null (unwatermarked) documents; TPR is measured at that threshold after the channel; ROC-AUC is computed from watermarked-vs-null score populations. Truthprint is run through its real linguistic codec on actual sentences ($\mu = 12$, $\tau = 20$, $[64, 32]$ GF(2) code); its quality is unigram token overlap (a BLEU-1-style precision without brevity penalty) between the watermarked and an independently realized copy of the same facts. Baseline fluency/quality is not measured because it requires each method's neural generator, absent from the Stage-1 abstraction.

Reading the tables. The clean setting saturates because the proxy generators deliberately select matching symbols and provide many carriers. Under the stipulated high-replacement setting, token proxies lose signal by construction while semantic proxies retain it according to the chosen erasure parameter. These outcomes validate the simulator implementation, not H1 and not any comparison among the named published systems. P2 and P3 separately exercise Truthprint's authenticated coding core (Section A); this harness

Table 4: Diagnostic clean-channel simulation (500 docs × 32 sentences; threshold at 1% FPR)†. Detection saturates for all methods on clean text; the discriminating transformed setting appears in Table 5.

Method	TPR@1%FPR	ROC-AUC	Overlap
KGW proxy	1.000	1.000	—
SemStamp proxy	1.000	1.000	—
SWAN proxy	1.000	1.000	—
Truthprint	1.000	1.000	0.84

†Controlled simulation only, not a neural or translation benchmark. Overlap (unigram token overlap with an independently realized copy) measured only for Truthprint (real sentences); baseline fluency needs each method's neural generator. Truthprint additionally applies a 20-bit tag check; this is not a measured end-to-end FPR.

Table 5: Diagnostic transformed-channel simulation (TPR at 1% FPR) $(\tau_{\text{tok}}, \epsilon_{\text{inv}})$ per language pair†: EN→KO $(0.92, 0.30)$, EN→HI $(0.95, 0.35)$, KO→EN $(0.90, 0.28)$. Outputs illustrate the stipulated channel equations, not real translation; rows are method proxies in the simulator, not implementations. Real-MT comparisons are Tables 1 and 12.

Method	EN→KO	EN→HI	KO→EN
KGW proxy	0.020	0.018	0.030
DEW proxy	0.000	0.000	0.000
SemStamp proxy	0.996	0.988	0.994
SWAN proxy	1.000	1.000	1.000
Truthprint	1.000	0.972	0.998

†Controlled Stage-1 channel; τ_{tok} (token replacement) and ϵ_{inv} (meaning-carrier erasure) per translation event. Only the Truthprint proxy in this harness exercises authenticated recovery; this table does not evaluate baseline security.

does not evaluate the security properties of the other systems.

G Reproducibility and Notation

This section pins every choice needed to reimplement the language-independent core exactly, so that the reference scripts and the paper agree symbol for symbol. Nothing here is left to interpretation; the two scripts in `reference/` realize precisely these definitions.

G.1 Primitive Instantiation

The reference implementation fixes the following concrete choices.

- **Invariant digest.** $h_I = \text{SHA-256}(\text{JSON}(I))$, where JSON is canonical: keys sorted, separators "`,`" and "`:`", UTF-8 encoded. Canonicalization makes h_I order- and whitespace-independent.

- **Authentication.** MAC = HMAC-SHA256; the tag is the first τ bits of the MAC over $m \parallel n \parallel h_I$. The payload is $p = m \parallel \text{tag}$ (length $|p| = \mu + \tau$).

- **Keyed carrier map (binary carriers).** $\pi_j = \text{LSB}\big(\text{HMAC-SHA256}(K, h_I \parallel n \parallel \text{ID}_j)\big)$ with ID_j the 4-byte big-endian carrier index. The realized option is $a_j^\star = \pi_j \oplus c_j$ and the recovered symbol is $\hat{c}_j = \pi_j \oplus a_j^{obs}$.

- **Error-correcting code.** A systematic $[N, k]$ linear code over GF(2) with generator $G = [\, I_k \mid P \,]$, where P is drawn once from a fixed pseudo-random seed. Erasure decoding solves the surviving linear equations by Gaussian elimination over GF(2); it succeeds iff the retained rows have rank k (Proposition 3).

- **Detection decision.** Attribute iff erasure decoding returns some $\hat{p}$ *and* MAC verifies on $(\hat{m}, n, h_I)$; otherwise report the calibrated statistic z of Eq. (2) against a threshold τ_z without asserting authenticated attribution.

Table 6: Mapping between paper notation and reference-implementation identifiers.

Symbol	Meaning	Code identifier
h_I	canonical invariant digest	`canonical_invariant_digest`
p	authenticated payload	`authenticate_payload`
MAC	tag verification	`verify_payload`
ECC	systematic GF(2) code	`GF2Code.encode`
decode	erasure decoding	`GF2Code.decode_erasure`
π_j	keyed carrier offset	`keyed_bit`
$a_j^\star$	realized option	`realize_option`
$\hat{c}_j$	recovered symbol	`recover_symbol`
D_K	detector	`detect_document`

G.2 Notation-to-Code Map

Table 6 maps paper symbols to identifiers in `reference/truthprint_poc.py`, removing any ambiguity between the formalism and the code.

G.3 Generator Plural Fix and Re-translation

An earlier generator pluralized by appending "s", so 153 of the 800 NLLB source sentences read *cache misss*. We fixed the generator, which changes exactly those 153 sources, and re-translated all 800 with the same NLLB system on CPU (`scripts/make_realmt.py`); all 3,882 translations of the 647 unchanged sources are identical to the original GPU run. Every real-MT number in this paper comes from the corrected run. The fix leaves core6 document attribution unchanged in all six conditions and moves sentence-level TPR by at most 0.015 (Hindi 0.886 → 0.901). Scoring the 153 sentences in both runs shows that the patient noun, not the misspelling, drives their behavior: Arabic still loses the patient on the corrected text (0.007 → 0.026), and core6 recovery is similar (0.837 → 0.843). These sentences are easier than average for core6 (which excludes the patient), so dropping them lowers Arabic document attribution (0.900 → 0.700 in both runs). The neural-defense run (Table 10) keeps the legacy text because its cached LLM outputs are keyed to it (`paper/results/plural_artifact.md`).

G.4 Models, Data, and Licenses

All evaluation text is synthetic and produced by the released code (MIT license). Translation uses NLLB-200-distilled-600M (NLLB Team et al., 2024) (CC-BY-NC-4.0, non-commercial research use); the token-watermark baselines generate with GPT-2 (MIT) and Qwen2.5-0.5B-Instruct (Apache-2.0), SynthID-Text through the `transformers` implementation (Apache-2.0); the encoders are `paraphrase-multilingual-MiniLM-L12-v2` (Apache-2.0) and `potion-multilingual-128M` (MIT); the neural parsers are GPT-4o-mini (OpenAI API terms) and Llama-3.3-70B-Instruct (Llama 3.3 Community License), queried through OpenAI-compatible APIs. Every use is research evaluation within these terms; because of the NLLB license, the measured pipeline is not itself a deployable commercial system. No model is trained. On four CPU threads (`torch` 2.14.1, `transformers` 5.20.2) rebuilding the 4,800 NLLB outputs takes about 30 minutes (`scripts/make_realmt.py`) and the same-text run about 1.5 hours (55 minutes of generation, 36 of translation); every scoring script finishes in minutes. The original NLLB and KGW runs (Kaggle GPU) and the neural-parser runs (hosted APIs) did not record wall-clock time.

G.5 Parameters, Determinism, and How to Run

The core script uses $\mu = 16$, $\tau = 32$, $[N, k] = [96, 48]$; the linguistic script uses $\mu = 12$, $\tau = 20$, $[N, k] = [64, 32]$ over 32 sentences with two carriers each. All randomness is seeded: message bits and erasure sampling use fixed seeds, and the code matrix P uses a fixed seed, so every reported number is reproducible. The document nonce n is freshly sampled per run and passed explicitly to the detector; recovery is invariant to its value, so success does not depend on the nonce. Both scripts return a nonzero exit code if any asserted property fails, so they double as regression tests. Reproduction requires only Python 3 (standard library only); the paper builds with `pdflatex` via `make`. The exact commands are:

```
python3 reference/truthprint_poc.py
        # P1-P4
python3 reference/truthprint_linguistic_demo.
    py # L1-L3
make
        # build main.pdf
```

Expected results are those reported in Section A: P1–P4 hold, and L1–L3 hold on sentence strings.

The same primitives are also packaged as an installable Python distribution (`code/` in the accompanying repository) with a unit-test suite and continuous integration. Its command-line interface reruns all properties (`truthprint selftest`) and regenerates Table 2 deterministically from a fixed seed with 500 trials per erasure rate (`truthprint repro-table`).

H Carrier Capacity and Minimum Reliable Length

This analysis concerns the in-text carrier channel (Secs. 7–8), which the body reports as a negative result under real MT (Sec. 10.1); it bounds when that channel can work at all.

Model each carrier j as a q_j-ary symbol on an erasure channel with erasure probability ϵ_j, where ϵ_j is the chance that a transformation

destroys or obscures that site. The expected recoverable information is

$$\mathbb{E}[\text{bits}] \approx \sum_j (1-\epsilon_j)\log_2 q_j. \tag{4}$$

With a rate-R code protecting an authenticated payload of $|p| = \mu + \tau$ bits, recovery succeeds with high probability once the surviving capacity exceeds $|p|$ with margin. For N uniform binary carriers ($q_j = 2$) with erasure rate ϵ and a systematic rate-R code ($k = RN$), this requires

$$(1-\epsilon)N \gtrsim k \iff \epsilon \lesssim 1-R. \tag{5}$$

Proposition 3 (Erasure recovery condition)**.** *For a fixed linear $[N, k]$ code, erasure decoding succeeds exactly when the generator-matrix rows corresponding to the surviving positions have rank k. Consequently at least k positions must survive, giving the necessary bound $\epsilon \leq 1 - R$ for a realized erasure pattern. Whether recovery occurs near that boundary depends on the code construction and erasure distribution; the reference code's finite-length transition is measured rather than claimed as a universal threshold.*

Equation (5) gives a necessary carrier-count condition $N \geq \lceil (\mu+\tau)/((1-\epsilon)\log_2 \bar{q}) \rceil$ with $R \leq 1-\epsilon$—a planning bound explaining why short or heavily summarized text falls below capacity (the finite-length transition in Appendix Table 2 is specific to the implemented code).

I Computational Complexity

Let a document have S discourse units, $|C|$ carriers, κ candidates per carrier, and N code positions. Encoding cost is $O(S \cdot |C| \cdot \kappa \cdot c_{\text{val}})$, dominated by the validator cost c_{val} (a parse plus entailment check). Detection cost is $O(S \cdot c_{\text{parse}} + N \cdot c_{\text{map}} + D_{\text{ECC}})$, where the keyed map c_{map} is one HMAC evaluation and $D_{\text{ECC}} = O(Nk)$ for GF(2) erasure elimination (worst case $O(Nk^2)$). The cryptographic and coding operations scale with the number of carriers and are independent of the generator parameter count. End-to-end latency is nevertheless unknown because candidate generation, reparsing, and validation dominate and may themselves invoke large models; it must be measured for each implementation.

J Encoding and Detection Pseudocode

Listing 1: Truthprint encoding pseudocode.

```
G <- BuildTruthprint(x)
I <- ExtractInvariants(G)
p <- AuthenticatePayload(K, m, nonce, Hash(I)
    )
c <- ErrorCorrectingEncode(p)
C <- DiscoverAdmissibleCarriers(G)
for assignment in KeyedAssign(K, I, C, c):
    G2 <- ApplyCarrier(G, assignment)
    candidates <- RealizeCandidates(G2,
    language)
    accepted <- ValidateAgainstContract(
    candidates, I)
    if accepted is empty:
        MarkErasure(assignment)
    else:
        CommitBestCandidate(accepted)
return AssembleDocument()
```

Listing 2: Truthprint detection pseudocode.

```
segments <- SegmentAndIdentifyLanguage(text)
observations <- []
for segment in segments:
    G, conf <- ParseTruthprint(segment)
    I <- CanonicalizeInvariants(G)
    for carrier in ExtractObservableCarriers
    (G):
        if Reliability(carrier, conf) >=
    threshold:
            observations.append(
    InvertKeyedMap(K, I, carrier))
        else:
            observations.append(ERASURE)
candidates <- ECCDecode(observations)
valid <- VerifyMACs(K, candidates)
score <- CalibratedStatisticalScore(
    observations, K)
return score, Best(valid), ConfidenceReport()
```

K Practical Software Architecture

A prototype can use the following modular layout:

```
truthprint/
  schema/        # entities, events,
    relations, mutability
  parser/        # multilingual parser and
    canonicalizer
  carriers/      # voice, packaging,
    discourse, ordering
  encoding/      # payload, keyed mapping,
    ECC
  realization/   # candidate generation and
    ranking
  validation/    # hard constraints, NLI,
    numbers, quality
  detection/     # extraction, decoding,
    statistics
  attacks/       # translation, paraphrase,
    summary, splice
```

```
  experiments/   # datasets, metrics,
    runners
```

K.1 Core Data Model

Listing 3: Illustrative Python data model.

```
from enum import Enum
from pydantic import BaseModel, Field

class Mutability(str, Enum):
    LOCKED = "locked"
    CONDITIONAL = "conditional"
    REALIZATION_ONLY = "realization_only"

class Event(BaseModel):
    id: str
    predicate: str
    roles: dict[str, str]
    polarity: str = "positive"
    modality: str = "asserted"
    time: dict = Field(default_factory=dict)
    quantity: dict | None = None
    attribution: str | None = None

class CarrierSite(BaseModel):
    carrier_id: str
    family: str
    target_node_ids: list[str]
    alternatives: list[str]
    recoverability: float
    quality_cost: float
```

K.2 Keyed Mapping Skeleton

Listing 4: Secret-keyed carrier assignment skeleton.

```
import hmac
from hashlib import sha256

def keyed_index(key, invariant_hash, nonce,
                carrier_id, symbol,
    option_count):
    msg = b"|".join([
        invariant_hash,
        nonce,
        carrier_id.encode(),
        symbol.to_bytes(4, "big")
    ])
    digest = hmac.new(key, msg, sha256).
    digest()
    return int.from_bytes(digest[:8], "big")
     % option_count
```

L Evaluation Coverage

What this paper did and did not evaluate, against the protocol of Section B. Items marked *not evaluated* are open, not implied.

L.1 Baselines

Evaluated: unwatermarked generation; KGW on GPT-2 (different source text) and on the same Qwen2.5-0.5B text; SynthID-Text (`transformers`, mean g-value) on the same text; our SemStamp-style re-implementation; embedding SimHash; MinHash; SHA-256; ledger retrieval alone and with an NLI proxy or oracle. *Not evaluated:* official SemStamp, SemaMark, semantic-invariant watermarking, SWAN, DEW, a trained SynthID detector, and a standardized toolkit such as MarkLLM (Pan et al., 2024).

L.2 Languages and Domains

Evaluated: English source with NLLB translation into Korean, Hindi, Chinese, Arabic, and German, plus a Korean round trip; one closed synthetic domain (software incidents). *Not evaluated:* news, encyclopedic, technical, scientific, policy, dialogue, and creative text, and low-capacity spans (code, formulas, legal clauses, quotations).

L.3 Attacks

Evaluated: direct and round-trip translation; rule-based benign paraphrase; single-field meaning-altering tampers; a closed-lexicon adaptive paraphrase and a held-out out-of-inventory attack (Section 10.7). *Not evaluated:* Unicode and spelling edits, two-hop translation, deletion, insertion, splicing, summarization, regeneration, and detector-guided rewriting.

L.4 Metrics

Reported: TPR at a matched 1% FPR, ROC-AUC (same-text), document attribution, sentence authentication, tamper rejection, per-field recovery, false attributions per ledger pair, and ValidRemoval. *Not reported:* TPR at 0.1/0.01% FPR, payload bit-error rate on real text, minimum reliable length on real text, latency, entailment-based integrity, and human factual-equivalence judgments.

The invariant error rate of Table 3 (not measured here) is

$$\mathrm{IER} = \frac{\sum_i w_i E_i}{\sum_i w_i}, \tag{6}$$

where critical categories such as polarity and numeric value receive larger weights or are treated as hard failures.

Table 7: Claim–evidence map for the present submission (part 1 of 2: method and robustness claims; part 2 is Table 8). "Not established" means that the claim must not be inferred from the diagnostic simulation.

Claim	Status	Supporting evidence or missing test
Coding/ authentication round trip	Demonstrated	Deterministic core tests P1–P4 and closed-domain string tests L1–L3.
Protected-field fidelity	Conditional (validator soundness measured)	Remark 1, assuming a sound invariant extractor and validator. Soundness is measured as the false-accept rate on meaning-altering edits on the closed domain (0 observed, C1); under real MT only a simulated tamper (one contract field of the registered record altered) is measured, failing to match at ≥ 0.996 (Sec. 10.4); wide-coverage soundness remains open.
Fresh-payload unforgeability	Conditional	Reduction to EUF-CMA security of the MAC; replay and nonce discovery are separate deployment concerns.
Surface-carrier translation robustness	Refuted (real MT)	Stage-1 surface parser reads 800/800 source sentences but 1/4,800 real NLLB translations (Sec. 10.1).
Meaning-layer recovery under real MT	Demonstrated (closed-domain, six languages; data committed)	Stage-2 extractor recovers fields at 0.83–0.99 over 4,800 real NLLB translations in six languages (closed categoricals ≥ 0.97; open entity/relational fields 0.83–0.92), and the meaning-digest authenticates documents cross-lingually (core6 document attribution 1.000 for German/Hindi, 0.90–0.98 for the rest; Secs. 10.2, 10.4). The full run ships in-repo (`paper/results/realmt_b2/`); wide-coverage parsing and larger document counts remain open.
Wide-coverage neural parsing on open-domain text	Pilot (positive per field; no sentence fully exact)	In a small LLM-MT pilot (16 base sentences, 6 with open wording; 96 records) a neural extractor beats the lexicon on open wording (0.614 vs. 0.481 field recovery) but recovers all nine fields exactly in none of the 96 records (lexicon 0.479) while the lexicon leads on closed templates; single-model and small-scale caveats apply (Sec. 10.3).
Benign paraphrase robustness	Demonstrated (closed-domain)	On 200 documents $\times 16$ sentences, meaning-preserving in-lexicon paraphrase (voice, time reflow, verb/time synonyms, discourse hedge) authenticates at TPR 1.000 and document attribution 1.000; single-field altering edits rejected 1.000 (Sec. 10.7).
Adaptive-attack robustness	Closed lexicon refuted; defended by extended coverage; open-vocabulary open	A schema-aware attacker swapping out-of-lexicon content synonyms ("took care of" for "fixed") reaches ValidRemoval 1.000 (Eq. 1) under the closed lexicon, localized to the stripped predicate/time/attribution fields (polarity and agent survive 100%). An extended-coverage frontend recovers those fields, driving adaptive ValidRemoval to 0.000 while benign auth (1.000) and tamper rejection (1.000) are unchanged; a held-out out-of-inventory attack still evades any fixed lexicon (ValidRemoval 1.000), but two small real open-vocabulary neural runs (GPT-4o-mini N=40, Llama-3.3-70B N=10) resist even the held-out attack (ValidRemoval 0.050/0.100, CIs disjoint from the lexicons') at benign/tamper 1.000. Coverage, not the meaning-digest, is the lever; the neural frontend supplies open-vocabulary coverage (Sec. 10.7, Table 10).

Table 8: Claim–evidence map (part 2 of 2, continued from Table 7): false-positive and baseline-comparison claims.

Claim	Status	Supporting evidence or missing test
System-level false-positive rate	Not established	The tag bound is per tested fresh candidate; parser search, tested key–nonce pairs, and fallback scores require end-to-end calibration.
Semantic-level false-positive rate (contract collision)	Measured (closed-domain)	The $2^{-\tau}$ bound covers cryptographic tag collision only; the meaning-digest floor is the contract-collision probability. On the closed-domain corpus this floor is 5.2×10^{-3} ($\approx 1/192$) for the six-field contract—matching the measured 0.000–0.007 real-MT false positives and $\sim 2 \times 10^{7}\times$ the 2^{-32} bound (Sec. 10.4). Open-domain text has lower contract entropy, so estimating it per corpus is a deployment prerequisite.
Comparative robustness vs. token and semantic baselines	Demonstrated, small scale (real MT: KGW token; matched-FPR SemStamp-style + robust hashing)	Under the same NLLB translator a KGW watermark on GPT-2 generations collapses (TPR@1%FPR 0.960 clean → 0.000–0.380 translated) while meaning-digest document attribution on the template corpus stays at 0.90–1.00; different source texts, so not a controlled head-to-head (Table 12, Sec. 10.5). At a matched 1% FPR on the 800-sentence NLLB corpus the typed digest identifies every translated document (1.000, six conditions) while a SemStamp-style re-implementation reaches 0.020–0.300 (templates handicap embedding methods; 0.377–1.000 on a separate 16-sentence set); with a paraphrase-tuned encoder SimHash also survives (0.990–1.000), so survival alone does not separate the methods, and only the typed contract rejects single-field tampers (Table 1, Sec. 10.5). On the same LLM text (Table 13) KGW and SynthID-Text fall to 0.000–0.133 after translation, Truthprint attributes 0.633–0.967 (30 documents, one 0.5B generator); SWAN and a trained SynthID detector remain future work.
Distinction from robust/semantic hashing	Demonstrated (closed-domain, real MT)	Robust hashes (exact, MinHash, SimHash) either die under translation or cannot detect a typed field change (reject ≤ 0.005 of single-field tampers vs. typed 1.000); the meaning-digest is a typed, keyed member of the family with an analyzed collision floor (Secs. 2, 10.5).
Distinction from store-and-retrieve (ledger) baseline	Demonstrated (closed-domain)	Under the same ledger, retrieval and a cheap NLI proxy reject $\leq$ 0.06 of single-field tampers at matched benign acceptance (≤ 0.20 with a real multilingual encoder); only a perfect (unrealizable) oracle NLI reaches 1.000, and needs a semantic model, which the typed digest does not (Sec. 10.6).
Translator-independent (separated) frontend	Pilot (open-domain positive on one field)	A model-separated embedding frontend recovers open-domain temporal direction (0.944) where the lexicon abstains (0.000); static embeddings remain insufficient for the full contract, motivating the separated-backend LLM frontend (Sec. 10.3).

Table 9: Positioning of Truthprint relative to representative watermarking and attack research.

Category	Representative work	Main carrier or focus	Limitation addressed by Truthprint
Token-level sampling	SynthID-Text (Dathathri et al., 2024)	Keyed token-selection statistics	Translation and paraphrasing replace token decisions.
Sentence-space watermarking	SemStamp (Hou et al., 2024)	Keyed sentence-embedding regions	Embedding regions are not typed by invariant category.
Semantic-context watermarking	SIR (Liu et al., 2024b), SemaMark (Ren et al., 2024)	Contextual semantic embeddings	Similarity does not explicitly protect negation, numbers, causality, or attribution.
Hybrid defense	SynGuard (Han et al., 2025)	Token watermark plus semantic alignment	Semantic constraints are less explicit than a typed IR contract.
Structured semantics	SWAN (Ye et al., 2026)	AMR-level structures	Existing semantic form is not a complete watermark-specific mutability contract.
Dual embeddings	DEW (Schäfer et al., 2026)	Context and token embeddings	Language invariance remains implicit rather than represented as a canonical IR.
Attack analysis	SynthID-Text analysis (Omidi et al., 2026), black-box probing (Gloaguen et al., 2025), adversarial paraphrasing (Cheng et al., 2025)	Detector assumptions and adaptive transformation	Motivates authenticated payloads, conservative interfaces, and erasure-aware decoding.
Proposed system	Truthprint	Typed invariants and controlled realization classes	Explicitly specifies what must remain fixed and what may carry symbols.

Table 10: Open-vocabulary neural frontend vs. fixed lexicons on the adaptive corpus (contract core6, cached real-LLM extractions; Wilson 95% CI). *Two* instruction LLMs resist the *held-out novel* attack that both fixed lexicons fail, at unchanged benign auth and tamper rejection. N = items scored per frontend; the adaptive-seen column scores three attack variants per item ($3N$). The corpus keeps the pre-fix generator text (Appendix G.3).

Frontend	Benign TPR	Adaptive-seen VR	Held-out novel VR	Tamper rej.
Closed lexicon (N=40)	1.000	1.000	1.000	1.000
Extended lexicon (N=40)	1.000	0.000	1.000	1.000
Neural: GPT-4o-mini (N=40)	1.000	0.050 $[0.023, 0.105]$	0.050 $[0.014, 0.165]$	1.000
Neural: Llama-3.3-70B (N=10)	1.000	0.200 $[0.095, 0.373]$	0.100 $[0.018, 0.404]$	1.000

Table 11: Matched-FPR comparison on a separate set of 16 base sentences translated by an LLM (no German condition); the 300 documents per condition resample these 16 sentences. Same methods and protocol as Table 1.

Method	RT	KO	HI	ZH	AR
Typed digest	1.000	1.000	1.000	1.000	1.000
SemStamp	1.000	0.937	0.617	0.460	0.377
SimHash	1.000	0.590	0.000	0.883	0.633
MinHash	1.000	0.000	0.000	0.000	0.000
SHA-256	0.867	0.000	0.000	0.000	0.000

Table 12: Token watermark vs. meaning digest under the same real NLLB translator. KGW (on GPT-2 generations) is scored per document at a calibrated 1% FPR; the Truthprint rows are core6 on the template-fact corpus at its contract-collision operating point ($\sim 5 \times 10^{-3}$). Source texts and operating points differ, so this shows token-signal removal, not a controlled head-to-head; the clean cell (—) is undefined for a ledger scheme with no in-text signal. Section 10.5 (Table 1) gives the matched-FPR comparison with a SemStamp-style detector and robust hashing. Replaces the diagnostic simulator (Appendix F).

Method	clean	RT	KO	HI	ZH	AR	DE
KGW (token)	0.960	0.380	0.040	0.040	0.020	0.000	0.080
Truthprint, doc. attribution	—	0.980	0.940	1.000	0.940	0.900	1.000
Truthprint, sentence TPR	—	0.720	0.684	0.901	0.666	0.619	0.850

Table 13: Same-text control: one generator (`Qwen2.5-0.5B-Instruct`, CPU) rewords the facts of 30 committed 16-sentence documents unwatermarked, with KGW (γ=0.25, δ=2), and with SynthID-Text (`transformers` tournament sampling, unweighted mean g-score); all are translated by NLLB-200-distilled-600M. Token watermarks: TPR above the maximum score of the 30 unwatermarked documents (empirical FPR 0/30) and ROC-AUC. Truthprint uses the unwatermarked text, accepting a sample only if Stage-2 reads back its core6 record (up to 4 samples; 0.871 [0.838, 0.898] of sentences validated, the rest count as failures), and attributes documents id-free against the 30 registered documents. False attributions: 0/30 in every condition. Wilson 95% CI in brackets. `paper/results/same_text_nllb.md`. Fact fidelity of the clean English text, by the extractor-based proxy (core6 read back; not human-judged): unwatermarked first sample 0.617 [0.572, 0.659], KGW 0.477 [0.433, 0.522], SynthID-Text 0.498 [0.453, 0.542]; the validation loop reaches 0.871 at a mean of 1.79 draws per sentence (`same_text_fidelity.md`).

Condition	KGW TPR	KGW AUC	SynthID TPR	SynthID AUC	TP doc. attr.	TP sent. auth.
clean	0.933 [0.787,0.982]	0.998	0.267 [0.142,0.444]	0.881	1.000 [0.886,1.000]	0.871
KO	0.100 [0.035,0.256]	0.609	0.000 [0.000,0.114]	0.672	0.633 [0.455,0.781]	0.508
AR	0.000 [0.000,0.114]	0.527	0.133 [0.053,0.297]	0.529	0.667 [0.488,0.808]	0.500
DE	0.033 [0.006,0.167]	0.631	0.000 [0.000,0.114]	0.499	0.967 [0.833,0.994]	0.731

Table 14: Ledger retrieval(+NLI) vs. the typed meaning-digest, same ledger and id-addressed verification, benign-paraphrase acceptance matched at ≈ 0.95 (960 closed-domain items, Wilson 95% CI). Storing and retrieving the source survives paraphrase but cannot reject a single-field tamper; only a perfect (unrealizable) NLI matches the typed contract.

Method	Benign TPR	Tamper rej.	Model
Meaning-digest (typed)	1.000	1.000	none
Retrieval (NN)	0.941	0.047	encoder
Retrieval + NLI (proxy)	0.918	0.061	NLI
Retrieval + NLI (oracle)	0.941	1.000	perfect NLI

Table 15: Paraphrase, adaptive attack, and extended-lexicon defense on real English strings (200 documents $\times$16 sentences, contract core6, Wilson 95% CI). Benign meaning-preserving paraphrase is authenticated (RQ4); a schema-aware adaptive attacker (RQ7) removes the mark under the *closed* lexicon (ValidRemoval 1.000), but an *extended-coverage* frontend recovers the same invariants and drives ValidRemoval to 0.000 while still rejecting altering edits; a held-out attack whose synonyms lie outside the extended inventory still evades (coverage bound). Auth TPR is over meaning-preserving paraphrases; tamper rejection for the altering row.

Transform family	Auth TPR	ValidRemoval	Surface carrier
Benign paraphrase	1.000	0.000	0.015
Adaptive (RQ7), closed lexicon	0.000	1.000	0.000
Adaptive (RQ7), extended	1.000	0.000	0.000
Held-out novel, extended	0.000	1.000	0.000
Altering (tamper), extended	0.000	— (1.000 rej.)	0.000